

DS4 Halo on AMD Strix Halo

DeepSeek V4 Flash 0731 IQ2 · Single Radeon 8060S · 128 GB RAM

Original DS4 · official published result

295.27

token/s · added 2K interval ending at 4K

Source: antirez/ds4 gfx1151 report

DS4 Halo · our recorded result

454.59

token/s · complete prepared 4K request

Samples: 454.52 / 454.65 token/s

Different benchmark protocols. No percentage speedup inferred between these rates.

Halo: 3 pure warmups + 1 measured request per process · IOMMU off · 106 GiB GPU ceiling

FP32 logits, complete serialized state and token IDs bitwise-identical in verified cases.

Model loading / generation excluded. No profiler, trace or payload readback during Halo timing.

Source, installer and evidence: github.com/MadRock-AI/ds4