


DS4 Halo on AMD Strix Halo

DeepSeek V4 Flash 0731 IQ2 · Single Radeon 8060S · 128 GB RAM

Our same-system comparison · Prepared complete 4K prefill · Two samples per engine



3 pure warmups + 1 measured request per process · IOMMU off · 106 GiB shared GPU limit

FP32 logits, complete serialized state and token IDs bitwise-identical in the verified cases.

Model loading / generation excluded. No profiler, trace or payload readback during timing.

Halo samples: 454.52 / 454.65 token/s. Source, installer and evidence: github.com/MadRock-AI/ds4