

Prefill on AMD Strix Halo


DeepSeek V4 Flash 0731 · Single device, 128 GB unified memory

Peak recorded mean: 449.03 token/s


Resident 4K after native preparation / warmup; separate from the first-use curves below.

Bitwise logits, complete state and tokens preserved in the verified cases.

Full prompt from empty context



Incremental prefill · 2K tokens added at each frontier



Model loading excluded; first-use preparation included. Incremental rates count only the added 2K tokens.

K = 1,024 tokens. Logarithmic context axis. Markers are archived measurements; lines only connect measured points.

Local upstream = rebuilt upstream8db control. Campaigns: September 25–26, 2026. No new benchmark.